



INTRO NATURAL L A N G U A G E P R O C E S S I N G (N L P)

– Matee Vadrukchid –



Big Picture of NLP

Big picture of NLP



1. Introduction and Outline

1.1. What We're Building

- A **mini conversation generator** that responds to simple inputs.
- Example conversation lines:
 1. **User:** "Hello, how are you?"
Bot: "I'm thank you and you?"
 2. **User:** "I'm good." **Bot:** "Great!"
- We'll store a small set of **20 or so** conversation lines (like QA pairs).
- We'll build a neural network model that, given an **input sentence**, generates the **output sentence** token by token, starting with a **special start token** and ending with a **special end token**.

1. Introduction and Outline

1.2. Lesson Plan

1. Theory

- Word → Vector (Vocabulary, tokens, embeddings)
- Sequence-to-sequence modeling
- Transformer basics

2. Dataset & Vocabulary

3. Tokenization & Special Tokens (`<sos>` , `<eos>` , `<pad>`)

4. Simple Transformer Model

5. Training

6. Inference (Generating new responses)

7. Assignments for deeper practice

2. Theory

2.1. Word → Vector (Embeddings)

- Why do we need word vectors?
 - Neural networks operate on numbers, not text.
 - We convert words to dense vectors (embeddings) so that semantically similar words have similar vector representations.
- Vocabulary
 - We create a set (or list) of known words from our conversation data (e.g., "hello", "how", "are", "you", "i'm", etc.).
 - Each unique word is given a token ID (e.g., word "hello" → token ID 4).

2. Theory

- Adding Special Tokens
 - <sos> : start of sentence.
 - <eos> : end of sentence.
 - <pad> : for padding to a fixed length.
 - <unk> : unknown words not in our vocab.

2. Theory


2.2. Sequence-to-Sequence (Seq2Seq)

- **Concept:** We feed an **input sentence** to an **encoder** which processes it into a hidden representation. A **decoder** then processes that representation (and previous tokens it generated) to **predict the next token**.
- **Key:** The decoder continues generating tokens until it outputs `<eos>`.

2. Theory

2.3. Transformer Basics

- Uses **attention** to see different parts of the sentence simultaneously.
- Works well for translations, conversation generation, etc.
- **Encoder**: reads input sequence.
- **Decoder**: generates output sequence, one token at a time, while looking at the encoder's hidden representation.
- We'll implement a **simplified** version for demonstration.



Link:<https://medium.com/@sayedebad.777/building-a-transformer-from-scratch-a-step-by-step-guide-a3df0aeb7c9a>