



INTRO REINFORCEMENT LEARNING (RL)

– Matee Vadrukchid –




Big Picture of RL



Introduction and Outline

What is Reinforcement Learning?


- In Reinforcement Learning, an agent interacts with an environment to learn how to perform a particular task well



Introduction and Outline

What is Reinforcement Learning?

- How is it different to the other learning paradigms?



- There is no supervisor, only a reward.
- The agent's actions affect the subsequent data it receives
- Feedback is delayed, and may be received after several actions

Introduction and Outline

Cat Agent

State: Sitting



Action: walk



Action: keep sitting



Stay hungry..

Observable
Come here!



Reward

Introduction and Outline

Examples of Reinforcement Learning


Fly a helicopter



Make a robot walk



Manage an investment portfolio



Play Atari games better
than humans


Introduction and Outline

Example of games


- Chess



- Go



- Carla



Explanation & Theory

Overview

- **GridWorld** is a simple, discrete environment where the world is represented as a grid. Each cell in the grid represents a state.
- **States and Actions:**
 - **States:** In our code, each cell is defined by its (row, col) coordinate. For example, in a 4×4 grid there are 16 states.
 - **Actions:** The agent can take one of four actions: Up, Down, Left, or Right.
- **Start and Goal:**
 - The **Start** state is fixed (top-left cell, i.e., `(0, 0)`).
 - The **Goal** state is set to the bottom-right cell `(rows - 1, cols - 1)`.
- **Obstacles:**
 - Certain cells (for instance, `(1, 1)`) are designated as obstacles.
 - If the agent tries to move into an obstacle, it receives a penalty (in this version, the episode terminates with a reward of -1).

Explanation & Theory

Out-of-Bound Actions

- If the agent attempts to move off the grid (e.g., moving up from the top row), the environment immediately gives a penalty of -1 and terminates the episode.
- This teaches the agent that certain actions are “illegal” and should be avoided.

Reward Structure

- **Goal:** Reaching the goal gives a reward of +1.
- **Invalid Actions:**
 - Moving out-of-bound or hitting an obstacle gives a reward of -1 and terminates the episode.
- **Neutral Moves:** Valid moves that do not result in hitting an obstacle or reaching the goal provide a reward of 0.

The Q-Learning Algorithm

Key Concepts

- **Q-Value:**
 - The Q-value $Q(s, a)$ estimates the expected cumulative reward the agent will receive by taking action a in state s and then following an optimal policy.
- **Q-Table:**
 - A table where each row corresponds to a state and each column corresponds to an action.
 - The goal of Q-learning is to learn this table by iteratively updating it.

The Q-Learning Algorithm

The Update Rule

The core update rule is:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

- α (learning rate) controls how much new information overrides old information.
- r is the immediate reward received after taking action a in state s .
- γ (discount factor) determines how much future rewards are valued relative to immediate rewards.
- $\max_{a'} Q(s', a')$ is the estimated maximum future reward from the next state s' .

The Q-Learning Algorithm

Epsilon-Greedy Exploration

- **Exploration vs. Exploitation:**
 - With probability ϵ , the agent chooses a random action (exploration) to gather more information about the environment.
 - With probability $1 - \epsilon$, the agent chooses the action with the highest Q-value (exploitation) to maximize its reward.
- **Epsilon Decay:**
 - Over time, ϵ decays, meaning the agent gradually shifts from exploration to exploitation as it becomes more confident in its learned Q-values.